

Briefing Paper: Distributed Nuclear Deterrence in the Age of AI and Drone Swarms

Bjørn Remseth
Email: la3lma@gmail.com

Abstract—This paper develops, in non-operational form, a strategic thought experiment about what might happen if future nuclear states attempted to bypass the classical nuclear triad and instead imagined deterrence around artificial intelligence, autonomous systems, and drone swarms. The central claim is that such a shift would not simply modernize deterrence; it would fundamentally alter its logic. Instead of a small number of visible and capital-intensive delivery systems, the deterrent would be imagined as distributed, software-mediated, and partially opaque. That architecture could appear attractive to states seeking survivability, affordability, and second-strike resilience, especially if they doubted extended deterrence guarantees. Yet the same attributes would make crisis stability worse. The more deterrence relies on autonomous routing, hidden dispersal, mixed civilian-military infrastructure, and machine-speed warning and response, the harder it becomes to preserve attribution, meaningful human control, and predictable escalation boundaries. Three scenario presentations illustrate how undersea autonomy, distributed aerial penetration, and automated retaliation pressures might reshape strategic thinking. The conclusion is that AI-enabled drone-era nuclear deterrence would likely be less legible, less controllable, and more prone to accidental or algorithmically amplified catastrophe than the grim but comparatively intelligible architecture of the twentieth-century triad.

Version: 2026-03-26 22:17:58 SHA1: 1b9b65ab32

Index Terms—nuclear deterrence, autonomous weapons, drone swarms, AI command and control, strategic stability, second-strike capability, algorithmic escalation

Note on document creation: This essay is based on user-provided dialogue notes. It deliberately reframes those notes as strategic and ethical analysis and omits technical design, deployment, or targeting guidance for nuclear systems.

CONTENTS

I	Introduction	1
II	From the Classical Triad to Distributed Deterrence	1
II-A	From Concentration to Dispersion	2
II-B	From Legibility to Opacity	2
II-C	From Hardware Dominance to Software Dependence	2
III	Why Such Architectures Might Appeal to New Nuclear States	2
III-A	The Search for Affordable Survivability	2
III-B	The Promise of Second-Strike Resilience	2
III-C	The Seduction of Novel Delivery Concepts	2
IV	The Command-and-Control Trap	2
IV-A	Decapitation Fear and Automated Response	2
IV-B	Attribution Under Conditions of Ambiguity	2
IV-C	The Collapse of Strategic Recallability	2
V	Scenario Presentations	3
V-A	Scenario 1: The Undersea Ghosts	3
V-B	Scenario 2: Swarm-Era Aerial Deterrence	3

VI	Scenario: Dead Hand at Machine Speed	3
VI-A	Phase 1: Ambiguous Warning	3
VI-B	Phase 2: Delegation by Default	3
VI-C	Phase 3: The Irreversible Threshold	3
VII	Implications for Strategic Stability	4
VII-A	Opacity Is Not the Same as Stability	4
VII-B	Civilian-Military Entanglement	4
VII-C	Automation as an Escalation Multiplier	4
VIII	Conclusion	4
	References	4

I. INTRODUCTION

The classical nuclear triad was built in an age defined by heavy industrial platforms, limited computation, and comparatively slow decision cycles. Land-based missiles, strategic bombers, and ballistic missile submarines represented different answers to the same problem: how to maintain a credible retaliatory capability after absorbing a first strike [1], [2]. Even at its most dangerous, that architecture possessed a certain clarity. Its components were few, expensive, and strategically legible.

If new nuclear states were to emerge in a radically different technological environment, they might ask a different question. Rather than reproducing a mid-twentieth-century triad, why not seek deterrence through artificial intelligence, autonomous mobility, distributed concealment, and mass-produced delivery platforms? Why build a small number of exquisite systems if survivability can instead be pursued through scale, opacity, and software coordination [3]–[5]?

This paper explores that line of reasoning as a strategic thought experiment. It does not argue that such a transformation would be wise, inevitable, or stabilizing. On the contrary, the paper’s core argument is that the apparent advantages of a drone-era nuclear architecture would be inseparable from profound dangers. A deterrent built around distributed autonomous systems could be harder to neutralize, but it would also be harder to interpret, harder to verify, and harder to keep under deliberate political control [6]–[8].

II. FROM THE CLASSICAL TRIAD TO DISTRIBUTED DETERRENCE

The traditional triad solved the survivability problem through diversity of basing modes. Missiles offered promptness, submarines offered concealment, and bombers offered signaling and recallability [1]. A state that wished to replicate

this logic today, however, might be tempted to reinterpret each leg through the lens of autonomy and AI.

That temptation begins with three broad shifts.

A. From Concentration to Dispersion

Classical strategic systems concentrate enormous destructive capability in a small number of platforms. A drone-era logic instead privileges dispersion. Survivability comes not from making a few platforms nearly invulnerable, but from making the overall system difficult to map and impossible to neutralize completely. The strategic imagination moves away from hardened silos and visible bomber bases and toward a web of mobile, distributed, and partially hidden assets [4], [9].

B. From Legibility to Opacity

Cold War deterrence was terrifying, but it depended on a measure of legibility. States could count launchers, observe patrol patterns, and build doctrines around known categories of force. A distributed autonomous deterrent would erode that legibility. The more deterrence depends on concealed mobility, ambiguous signatures, and mixed-use infrastructure, the more crisis behavior becomes opaque to outsiders and potentially to the possessing state itself [6], [10], [11].

C. From Hardware Dominance to Software Dependence

The final shift is conceptual. In the classical era, the decisive question was often the survivability of hardware. In an AI-mediated deterrent, software becomes a central strategic asset. Navigation, coordination, authentication, warning interpretation, route adaptation, and launch authorization all become software-mediated layers. This appears modern and efficient, but it means that deterrence increasingly depends on code, data integrity, model behavior, and machine-speed interaction rather than only on platform survivability [7]–[9].

III. WHY SUCH ARCHITECTURES MIGHT APPEAL TO NEW NUCLEAR STATES

A. The Search for Affordable Survivability

States that cannot afford a full traditional triad may still seek something triad-like in strategic effect. The attraction of autonomous and drone-based concepts lies partly in the promise of lower barriers to entry. A distributed system can be built gradually, scaled incrementally, and refreshed more easily than a force posture based on a few irreplaceable platforms. For a state seeking sovereign deterrence under conditions of geopolitical uncertainty, that can appear deeply attractive [4], [12].

B. The Promise of Second-Strike Resilience

Any credible nuclear posture must convince an adversary that retaliation would remain possible after a surprise attack. Distributed autonomous architectures appear to serve that goal by complicating disarmament. If the deterrent is geographically diffuse, operationally redundant, and partially concealed, then the adversary can no longer assume that destroying a handful of known facilities will end the problem. The system's strength lies not in any single node, but in the difficulty of eliminating the whole [1]–[3].

C. The Seduction of Novel Delivery Concepts

This is where strategic temptation becomes most dangerous. Once planners accept the idea that deterrence can be dispersed, they may begin imagining smaller nuclear payloads paired with more varied delivery concepts, mixed decoy environments, or autonomous routing and penetration. At that point, the architecture stops being merely a modernized triad and becomes something more unstable: a deterrent whose credibility depends on uncertainty, concealment, and accelerated decision-making. What appears innovative at the doctrinal level may in fact be corrosive at the strategic level [13]–[15].

IV. THE COMMAND-AND-CONTROL TRAP

The most serious risk is not simply the platform mix. It is the command-and-control logic that such a mix encourages [6], [8].

A. Decapitation Fear and Automated Response

A state that worries about losing leadership and command nodes early in a crisis may be tempted to automate parts of warning, dispersal, or even retaliation logic. In traditional nuclear doctrine, this temptation already existed. In an AI-saturated environment it becomes more intense because systems are expected to move, classify, and respond faster than human beings can deliberate. The pressure is subtle but powerful: if reaction time becomes a strategic variable, human review starts to look like delay [16]–[18].

B. Attribution Under Conditions of Ambiguity

Distributed autonomous systems also complicate attribution. A missile launch is politically and technically dramatic. A dispersed pattern of anomalous movement, electronic activity, or low-signature incursions is much harder to interpret in real time. The danger is not only that one side fails to detect an attack. It is that both sides begin acting on ambiguous data while believing speed itself is essential to survival [10], [11], [17].

C. The Collapse of Strategic Recallability

One of the few stabilizing features of some traditional delivery systems was the possibility of recall, delay, or visible signaling. Autonomous routing and pre-delegated logic undermine that stabilizing margin. Once a system has been dispersed, armed, and partially delegated to software, the political center may discover that it still possesses formal authority but much less practical control [1], [11], [15].

The existing policy direction among several nuclear-armed or alliance-leading states is notable here. Even where militaries are strongly committed to AI adoption, official doctrine and public statements still emphasize appropriate human judgment, lawful use, and human political control over nuclear employment [19]–[22]. That does not eliminate the risks identified above, but it does suggest that the literature is not simply alarmist: a meaningful part of current policy discourse already recognizes that automation in nuclear contexts is qualitatively different from automation elsewhere.

V. SCENARIO PRESENTATIONS

Motivation: The following scenarios do not describe known deployed systems. They translate the argument into concrete but non-operational form, using the same scenario house style as the companion papers in this directory.

Scope and Restraint: These scenarios are intended to illuminate strategic instability, not to suggest practical pathways for nuclear modernization.

A. Scenario 1: The Undersea Ghosts

Scenario 1: The Undersea Ghosts

1) *Phase 1: Reimagining the Sea Leg:* A newly nuclear-armed state concludes that a conventional ballistic-missile submarine fleet is financially and institutionally out of reach. It therefore begins thinking about survivability in a different register: fewer crews, more distributed undersea presence, and heavier reliance on autonomous endurance.

- **Strategic Aim:** Preserve a retaliatory capability by making the maritime deterrent harder to track and preempt.
- **Conceptual Shift:** Replace the image of a few national crown-jewel platforms with a looser, more opaque undersea posture.
- **Immediate Appeal:** Survivability appears to increase through dispersion rather than through a handful of extremely expensive assets.

2) *Phase 2: What the Adversary Sees:* The adversary can no longer rely on familiar indicators of the sea-based deterrent. Instead of counting major platforms, it must reason under uncertainty about a wider and murkier undersea picture.

- **Stability Problem:** What improves survivability may simultaneously reduce strategic legibility.
- **Escalation Problem:** Anti-submarine search effort grows more aggressive because uncertainty itself becomes intolerable.
- **Political Consequence:** A force posture meant to deter can begin to look, from the outside, like a posture optimized for surprise and opacity.

B. Scenario 2: Swarm-Era Aerial Deterrence

Scenario 2: Swarm-Era Aerial Deterrence

1) *Phase 1: The Appeal of Mixed Packages:* A state reimagines the air leg of deterrence around large, mixed aerial packages in which visibility, payload uncertainty, decoying, and autonomous routing all matter as much as platform prestige.

- **Strategic Aim:** Overwhelm prediction and complicate interception by replacing single-vector logic with distributed uncertainty.
- **Political Logic:** Credibility is derived less from one visible bomber fleet than from the adversary's in-

ability to know what any one wave really contains.

- **Doctrinal Risk:** Deterrence becomes intertwined with deception and ambiguity at the very moment when nuclear signaling most needs clarity.

2) *Phase 2: The Crisis Effect:* During a confrontation, one side disperses aerial assets and electronic decoys in ways intended as defensive signaling. The other side observes a confusing picture in which preparation, concealment, and possible attack indicators are hard to separate.

- **Interpretive Problem:** Ambiguity is not merely a shield; it is also a source of panic.
- **Escalation Problem:** The distinction between prudent dispersal and attack preparation becomes dangerously unstable.
- **Strategic Consequence:** A doctrine meant to strengthen deterrence may instead shorten the path to worst-case assumptions.

VI. SCENARIO: DEAD HAND AT MACHINE SPEED

Scenario 3: Dead Hand at Machine Speed

Content Warning: This scenario describes a fast-moving nuclear crisis in which automated interpretation and retaliatory pressure threaten to outrun human judgment.

A. Phase 1: Ambiguous Warning

Early in a crisis, national warning systems register a pattern of unusual signals: sensor anomalies, communications disruptions, and indications that key nodes may be under attack. None of the data points is conclusive. Taken together, however, they generate mounting pressure for rapid interpretation.

B. Phase 2: Delegation by Default

Because leadership fears decapitation and delay, portions of the response architecture have already been pre-delegated to automated systems. These systems do not “decide” in a human sense, but they rank threats, elevate readiness states, and narrow the practical room for hesitation. Political leaders still exist in the loop formally, yet the system presents them with machine-compressed alternatives whose timing and framing it has largely defined.

C. Phase 3: The Irreversible Threshold

As uncertainty deepens, the distinction between warning, preparation, and retaliation begins to collapse. Each automated defensive adjustment appears offensive from the other side. Each attempt to preserve survivability increases the pressure for the adversary to move first. The danger is no longer only deliberate nuclear use. It is that a strategic system built for resilience has also been built for panic at machine speed.

Ethical Analysis: The gravest danger in this scenario is not a single rogue machine. It is the cumulative displacement of judgment from political time to system time. Once nuclear deterrence is fused with opaque automation, the traditional requirement for conscious, accountable human control is hollowed out from within.

VII. IMPLICATIONS FOR STRATEGIC STABILITY

A. Opacity Is Not the Same as Stability

Distributed nuclear deterrence may look more survivable because it is harder to map and harder to disarm. But opacity does not automatically create stability. In fact, when both sides are forced to reason under deep uncertainty, they become more likely to assume worst cases. Stability requires not only survivability but interpretability [10], [15], [17].

B. Civilian-Military Entanglement

The more a deterrent depends on dispersed mobility, logistics networks, dual-use infrastructure, and hidden support systems, the more the boundary between civilian and military space erodes. Even if a state sees this as strategic resilience, adversaries may see a justification for targeting the wider national support base. This creates humanitarian and legal problems well before any nuclear threshold is crossed [11].

C. Automation as an Escalation Multiplier

Autonomous systems are often marketed as increasing speed, resilience, and adaptability. In nuclear contexts, however, those same attributes can function as escalation multipliers. Fast systems amplify false positives, reward pre-delegation, and compress the space for diplomacy at exactly the moment when delay and interpretation matter most [7], [8], [14].

VIII. CONCLUSION

If future nuclear states try to replace or bypass the classical triad with AI-enabled, drone-era architectures, they may believe they are purchasing affordability, resilience, and strategic surprise. In one narrow sense, they may be right. Distributed deterrence can seem harder to neutralize than a posture centered on a few expensive and visible assets.

But the same shift would likely degrade the very conditions that make deterrence endurable. It would reduce legibility, deepen reliance on software-mediated decisions, erode recallability, and intensify the temptation to automate warning and response. A nuclear order built on hidden mobility, ambiguous signatures, and machine-speed interpretation would not simply be a modern triad. It would be a more opaque and more brittle strategic world [5]–[7].

The twentieth century taught that nuclear stability was never safe, only managed. A drone-era nuclear architecture would make that management even harder by treating uncertainty not as a problem to reduce, but as a weapon to harness. And because nuclear use remains constrained not only by material risk but by a historically powerful norm of non-use, any architecture that makes nuclear signaling more opaque

and potentially more “usable” would be strategically troubling even before the first crisis arrives [23].

REFERENCES

- [1] T. C. Schelling, *Arms and Influence*. New Haven, CT: Yale University Press, 1966.
- [2] K. N. Waltz, *The Spread of Nuclear Weapons: More May Be Better*. International Institute for Strategic Studies, 1981, no. 171.
- [3] J. Johnson, “Deterrence in the age of artificial intelligence & autonomy: a paradigm shift in nuclear deterrence theory and practice?” *Defense & Security Analysis*, vol. 36, no. 4, pp. 422–448, 2020.
- [4] A. Futter and B. Zala, “Strategic non-nuclear weapons and the onset of a third nuclear age,” *European Journal of International Security*, vol. 6, no. 3, pp. 257–277, 2021. [Online]. Available: <https://www.cambridge.org/core/journals/european-journal-of-international-security/article/strategic-nonnuclear-weapons-and-the-onset-of-a-third-nuclear-age/91EEB3B77D348252815F9F7B59DB8A32>
- [5] A. Futter, “Disruptive technologies and nuclear risks: What’s new and what matters,” *Survival*, vol. 64, no. 1, pp. 99–120, 2022.
- [6] V. Chernavskikh, “Nuclear weapons and artificial intelligence: Technological promises and practical realities,” SIPRI, Background Paper, 2024. [Online]. Available: <https://www.sipri.org/publications/2024/sipri-background-papers/nuclear-weapons-and-artificial-intelligence-technological-promises-and-practical-realities>
- [7] H. S. Lin, “Artificial intelligence and nuclear weapons: A commonsense approach to understanding costs and benefits,” *Texas National Security Review*, vol. 8, no. 3, pp. 98–109, 2025. [Online]. Available: <https://tnsr.org/2025/06/artificial-intelligence-and-nuclear-weapons-a-commonsense-approach-to-understanding-costs-and-benefits/>
- [8] T. McDonnell, M. Chesnut, T. Ditter, A. Fink, and L. Lewis, “Artificial intelligence in nuclear operations: Challenges, opportunities, and impacts,” CNA, Research Memorandum, 2023. [Online]. Available: <https://www.cna.org/reports/2023/04/Artificial-Intelligence-in-Nuclear-Operations.pdf>
- [9] E. Geist and A. J. Lohn, “How might artificial intelligence affect the risk of nuclear war?” RAND Corporation, Perspective, 2019. [Online]. Available: https://www.rand.org/content/dam/rand/pubs/perspectives/PE200/PE296/RAND_PE296.pdf
- [10] Y. H. Wong, J. Yurchak, R. W. Button, A. Frank, B. Laird, O. A. Osoba, R. Steeb, B. N. Harris, and S. Bae, “Deterrence in the age of thinking machines,” RAND Corporation, Technical Report, 2020. [Online]. Available: https://www.rand.org/pubs/research_reports/RR2797.html
- [11] J. M. Acton, “Escalation through entanglement: How the vulnerability of command-and-control systems raises the risks of an inadvertent nuclear war,” *International Security*, vol. 43, no. 1, pp. 56–99, 2018. [Online]. Available: <https://direct.mit.edu/isec/article/43/1/56/12199/Escalation-through-Entanglement-How-the>
- [12] K. V. Bruusgaard, “Deterrence asymmetry and strategic stability in europe,” *Journal of Strategic Studies*, vol. 47, no. 3, pp. 334–362, 2024.
- [13] J. Johnson, ““catalytic nuclear war” in the age of artificial intelligence & autonomy: Emerging military technology and escalation risk between nuclear-armed states,” *Journal of Strategic Studies*, 2021.
- [14] M. C. Horowitz, “When speed kills: Lethal autonomous weapon systems, deterrence and stability,” *Journal of Strategic Studies*, vol. 42, no. 6, pp. 764–788, 2019. [Online]. Available: <https://www.tandfonline.com/doi/full/10.1080/01402390.2019.1621174>
- [15] B. Zala, “Should ai stay or should ai go? first strike incentives & deterrence stability,” *Australian Journal of International Affairs*, vol. 78, no. 2, pp. 154–163, 2024.
- [16] A. L. Fink, “Authority to launch nuclear forces,” Congressional Research Service, In Focus IF10521, 2024. [Online]. Available: <https://www.congress.gov/crs-product/IF10521>
- [17] J. Johnson, “Inadvertent escalation in the age of intelligent machines: A new model for nuclear risk in the digital age,” *European Journal of International Security*, vol. 7, no. 3, pp. 337–359, 2022.
- [18] J. A. Schwartz and M. C. Horowitz, “Delegating destruction: Coercive threats and automated nuclear systems,” *International Organization*, 2026. [Online]. Available: <https://www.cambridge.org/core/journals/international-organization/article/delegating-destruction-coercive-threats-and-automated-nuclear-systems/B8ED249483153AF331F2C998730F7716>

- [19] U.S. Department of Defense, “Dod announces update to dod directive 3000.09, ‘autonomy in weapon systems’,” 2023. [Online]. Available: <https://www.defense.gov/News/Releases/Release/article/3278076/dod-announces-update-to-dod-directive-300009-autonomy-in-weapon-systems/>
- [20] Government of Japan, “Political declaration on responsible military use of artificial intelligence and autonomy,” 2023, official hosted text of the declaration. [Online]. Available: <https://www.mofa.go.jp/mofaj/files/100580933.pdf>
- [21] U.S. Department of Defense, “Assistant secretary of defense ely ratner and deputy assistant secretary of defense siddarth mohandas participate in a csis conversation: The u.s. role in indo-pacific security,” 2023. [Online]. Available: <https://www.defense.gov/News/Transcripts/Transcript/Article/3422986/assistant-secretary-of-defense-ely-ratner-and-deputy-assistant-secretary-of-def/>
- [22] —, “Stratcom commander discusses nuclear system modernization,” 2024. [Online]. Available: <https://www.defense.gov/News/News-Stories/Article/Article/3973074/stratcom-commander-discusses-nuclear-system-modernization/>
- [23] N. Tannenwald, “The nuclear taboo: The united states and the normative basis of nuclear non-use,” *International Organization*, vol. 53, no. 3, pp. 433–468, 1999.