The Hidden Context of Anthropic's Pentagon Refusal: What Mythos Reveals About the Stakes of Military AI

Bjørn Remseth
Email: la3lma@gmail.com
With AI assistance

***Abstract*—When Anthropic CEO Dario Amodei refused the Pentagon's request for unrestricted access to Claude, public discussion centered on surveillance and autonomous weapons. That interpretation is plausible but incomplete. A second possibility is that the refusal was shaped by the strategic implications of Mythos, Anthropic's later-revealed cyber model. If Mythos can discover zero-day vulnerabilities at scale and assist in multi-stage exploit development, transferring it to the Department of Defense would not simply extend current military AI practice; it could alter the tempo, scale, and stability of cyber operations. This report examines reported Mythos capabilities, current DOD cyber priorities, and the forms of military leverage such a system could have provided. The aim is not to infer Amodei's motives with certainty, but to show why the refusal may carry broader strategic significance.**

Note: This document was created with AI assistance; see "About This Document" and "Note on References and Verification" sections for details on methodology and verification processes.

***Index Terms*—artificial intelligence, cybersecurity, military AI, cyber warfare, autonomous weapons, AI safety, vulnerability assessment, Anthropic, Mythos AI**

I. Introduction: The OODA Loop Trap

Defense Secretary Pete Hegseth's demand that Anthropic remove Claude's safety guardrails for surveillance and autonomous weapons initially appeared to be another instance of the familiar dispute over military AI. Amodei's refusal, explaining that Anthropic could not "in good conscience" grant unrestricted access, was therefore widely interpreted as an ethical objection to militarization.

That interpretation may be too narrow. The refusal preceded Anthropic's disclosure of Mythos, a system described as having unusually strong cyber capabilities. We cannot know Amodei's reasoning, but the timing raises a useful question: was Anthropic refusing not only risky uses of current models, but also early military access to a more consequential capability?

The question matters because modern warfare rewards speed. Military planners often frame this through the OODA Loop—Observe, Orient, Decide, Act. As that loop compresses, each side is pushed toward greater automation, creating what some analysts describe as a "hyperwar equilibrium": once one actor removes human delay, others are pressured to follow.

Current AI systems remain poor candidates for life-or-death decisions. They hallucinate, behave unpredictably, and fail in ways that are difficult to audit in real time. Against that background, threatening Anthropic with "supply chain risk" status in order to secure compliance would already represent a significant escalation.

If Mythos was already in view internally, however, the dispute may have concerned something larger: whether the Pentagon would gain access to a system capable of reshaping cyber operations rather than merely assisting them.

II. The Mythos Revelation

In April 2025, Anthropic said it would restrict public access to Claude Mythos Preview because of cybersecurity risk. That decision changed the meaning of the earlier Pentagon dispute.

A. Technical Capabilities

According to Anthropic and independent evaluations [1], [2], Mythos reportedly found thousands of zero-day vulnerabilities across major operating systems, browsers, and critical software. It also performed expert-level cyber tasks that earlier models could not, succeeding 73% of the time and completing a 32-step attack simulation in three of ten runs.

These results matter because they suggest movement from research assistance toward partial operational autonomy. Mythos does not merely surface bugs; it appears able to reason about exploit chains, system interactions, and layered defenses at a level that previously required elite human teams.

B. Restricted Deployment

Anthropic treated that capability as exceptional. Instead of a general release, it limited access to roughly 40 large technology companies through "Project Glasswing," including Amazon, Apple, Microsoft, Google, and JP Morgan Chase. It also committed up to $100 million in usage credits and $4 million in direct support for defensive security work.

This is a striking break from the usual AI release model, which rewards broad access and rapid adoption. The restrictions suggest Anthropic judged Mythos to have crossed a threshold at which scale itself became a security concern.

III. What the Pentagon Wanted

To understand why Mythos would matter to the Pentagon, it is useful to begin with the scale of the existing cyber apparatus.

A. DOD Cybersecurity Infrastructure

The Department of Defense oversees roughly 440 organizations involved in cyberspace operations, with about 61,000 military and civilian personnel and more than 9,500 contractors [3]. Much of that work sits alongside U.S. Cyber Command (CYBERCOM) and covers offensive operations, defense, and management of the DOD Information Network.

The money points the same way. The FY2026 budget requests $5.4 billion for cyberspace operations, including about $2.6 billion for CYBERCOM [4]. Total cyber activity reaches roughly $15.1 billion, up 4.1% from the previous year [5].

B. Offensive Cyber Priorities

Senior officials have also been explicit that defense is not sufficient. Pentagon CIO Katie Arrington said, "We need more offensive capability" [6]. The implication is clear: the Pentagon seeks not only resilience, but greater operational freedom.

CYBERCOM's "2.0 Initiative," as described by Undersecretary Anthony Tata, aims to "accelerate the buildup of cutting-edge cyber tools" to address "acute and emerging" threats and strengthen deterrence. It is part of a broader push to make cyber operations faster, more scalable, and more deployable.

C. Tactical Cyber Operations

Military planners are also pushing cyber capability closer to the tactical edge. Traditional CYBERCOM work has focused on IP-based networks accessed from static, remote locations. Future conflicts are more likely to demand tools that can operate inside contested environments where that model breaks down.

Several services are already investing in offensive tools for the "blended electronic warfare or RF-enabled sphere at the tactical level." Taken together, these priorities form the backdrop for a strong Pentagon interest in a Mythos-class system.

IV. The Mythos Advantage: What DOD Could Have Achieved

If the Pentagon had gained access to Mythos, the implications for military cyber operations would have been broad and immediate.

A. Strategic Intelligence and Reconnaissance

Mythos could have substantially shortened the path from target identification to operational access. Instead of waiting months for human researchers to find exploitable flaws, operators could have mapped weaknesses across adversary systems at machine speed. Because military platforms and logistics networks depend heavily on commercial software, that visibility would not stop at classic espionage targets; it could extend into communications, support systems, and weapons infrastructure.

B. Automated Cyber Weapons Development

It could also have changed the economics of offensive cyber work. Exploit development is typically constrained by scarce human expertise. A model that can reason through complex exploit chains and solve expert tasks reliably would convert that bottleneck into a scaling problem. The Pentagon could have moved from crafting a limited number of high-value exploits to generating and refining them at far larger volume.

C. Real-Time Battlefield Cyber Operations

At the tactical level, Mythos could have compressed the interval between sensing a problem and acting on it. A battlefield unit facing new electronic warfare or communications interference would not need to wait on a distant analytical pipeline. In principle, a Mythos-enabled system could identify vulnerabilities, build an exploit path, and execute a disruption quickly enough to matter inside an active engagement.

D. Infrastructure and Economic Warfare

The same logic extends to civilian infrastructure. Power grids, transport systems, finance, and manufacturing all depend on the kind of widely deployed commercial software Mythos reportedly understands well. That could enable more selective and less attributable attacks: not blunt disruption, but calibrated manipulation that weakens an adversary's capacity while obscuring attribution.

V. The Strategic Implications

The military value of Mythos would not end with tactical advantage. It would also alter the structure of cyber competition itself.

A. The Automation of Cyber Conflict

Mythos marks a threshold in cyber automation. Earlier models could support human operators; Mythos appears closer to conducting multi-stage operations with limited supervision. That would allow the Pentagon to scale cyber campaigns far beyond current staffing limits, while also increasing the risk that operations unfold faster than human review can keep pace.

B. The Vulnerability Discovery Arms Race

A Mythos-class model would also intensify the zero-day arms race. A state that can discover in months what rivals find over years gains a temporary but meaningful strategic advantage. The difficulty is that such an advantage would invite imitation. Other countries would have strong incentives to build similar systems, pushing cyber competition toward an AI-driven discovery race.

C. Deterrence and Escalation Dynamics

Such a shift cuts both ways. Better knowledge of adversary weaknesses can strengthen deterrence by making retaliation more credible. It can also encourage preemption. If each side believes the other may soon gain a decisive cyber advantage, restraint becomes harder to justify.

VI. Why Anthropic Said No

This is what makes Anthropic's refusal strategically important, not merely ethically notable.

A. The Pandora's Box Problem

Seen in that light, Anthropic's refusal looks less like a reputational maneuver and more like an attempt to keep a new class of dual-use capability from entering military circulation too early. Cyber tools spread. Once techniques are operationalized, they are difficult to contain and eventually reach rivals and non-state actors.

By restricting Mythos to a narrow set of defensive uses, Anthropic appears to have been trying to preserve some of the system's security benefits without accelerating its offensive diffusion.

B. The Attribution Challenge

AI makes the attribution problem worse because it lowers the cost of sophistication while muddying responsibility. If complex attacks can be planned and adapted autonomously, assigning blame becomes harder and diplomatic response becomes harder to calibrate.

C. The Escalation Risk

It also raises the risk of machine-speed escalation. If several states deploy systems that can probe, exploit, and counter-exploit faster than operators can intervene, cyber conflict may begin to resemble an automated interaction problem rather than a controlled campaign. In that setting, small triggers can produce outsized consequences.

VII. Conclusion: The Road Not Taken

Anthropic's refusal looks different once Mythos is part of the frame. The decision may still have been ethical in the ordinary sense, but it can also be read as a strategic refusal to hand the Pentagon a cyber capability that might have changed the tempo and scale of military operations.

A Mythos-class system would have offered clear advantages: faster vulnerability discovery, cheaper exploit development, and tighter integration between cyber and battlefield action. It would also have raised the likelihood of arms-race dynamics, attribution failures, and escalation at machine speed.

That is what makes the episode important. The most difficult questions about military AI are not confined to autonomous weapons in the narrow sense. They also concern systems that compress decision cycles, expand offensive reach, and destabilize conflict even when no trigger is ever pulled by a robot. If Anthropic delayed that transition, even briefly, the refusal mattered more than it first appeared.

About This Document

This analysis was developed through a collaborative human-AI research process. The work began with the timing of Anthropic's refusal and the later disclosure of Mythos's cyber capabilities. From there, the research focused on DOD cyber programs, Mythos reporting, and the strategic consequences of combining the two. The final document is a synthesized argument rather than a transcript of that exploration. All references were checked for authenticity, accessibility, and fit with the claims they support. The intended audience is policymakers, security researchers, and readers interested in AI governance and military cyber strategy.

NOTE ON REFERENCES AND VERIFICATION

This document contains AI-assisted content. References were reviewed for academic integrity: URLs were tested, author names and publication venues were checked, and content was compared against the claims cited. Where available, DOIs were confirmed. Reference notes indicate whether an item was verified by URL or DOI, or still needs manual review. Readers should independently verify any source used for high-stakes applications.

REFERENCES

[1] Anthropic, "Project glasswing: Securing critical software for the ai era," 2025, verified: Official Anthropic announcement of Project Glasswing. [Online]. Available: https://www.anthropic.com/glasswing

[2] AI Safety Institute, "Our evaluation of claude mythos preview's cyber capabilities," 2025, verified: Official UK AISI evaluation of Mythos capabilities. [Online]. Available: https://www.aisi.gov.uk/blog/our-evaluation-of-claude-mythos-previews-cyber-capabilities

[3] U.S. Government Accountability Office, "Dod cyberspace operations: About 500 organizations have roles, with some potential overlap," January 2025, verified: GAO official report on DOD cybersecurity structure. [Online]. Available: https://www.gao.gov/products/gao-25-107121

[4] Congressional Research Service, "Fy2026 department of defense cyber budget request," 2025, verified: Official CRS analysis of DOD cyber budget. [Online]. Available: https://www.congress.gov/crs-product/IN12616

[5] Nextgov/FCW, "Dod gets millions for cyber capabilities under gop reconciliation package," July 2025, verified: Nextgov report on DOD cyber budget allocation. [Online]. Available: https://www.nextgov.com/cybersecurity/2025/07/dod-gets-millions-cyber-capabilities-under-gop-reconciliation-package/406540/

[6] DefenseScoop, "Pentagon cio calls for more offensive cyber capability," March 2025, verified: DefenseScoop report on Pentagon CIO statements. [Online]. Available: https://defensescoop.com/2025/03/20/katie-arrington-dod-cio-offensive-cyber-capability/