

Briefing Paper: Drone Swarms, Deterrence, and the Prospect of Non-Nuclear Mass Disruption

Bjørn Remseth
Email: la3lma@gmail.com

Abstract—This briefing paper transforms a dialogue-based thought experiment into a structured essay on how drone warfare could alter deterrence logic in a Europe marked by uncertainty about alliance guarantees. The core claim is not that drones become a direct substitute for nuclear weapons, but that large-scale autonomous swarms could emerge as instruments of non-nuclear mass disruption. Their strategic effect would come less from blast than from the rapid paralysis of power grids, communications, logistics, and public administration. In that sense, the relevant comparison is not only with weapons of mass destruction, but also with systems capable of producing nationwide societal breakdown without crossing the nuclear threshold. The paper examines why such systems might appeal to states seeking sovereign deterrence, why they would also destabilize crisis management, and how they could compress military and political decision cycles to machine speed. Three scenario presentations adapt the house style used elsewhere in this directory to show how distributed drone deterrence, infrastructure blackout attacks, and AI-mediated border crises might unfold. The conclusion is sobering: drone-based deterrence promises affordability, deniability, and scale, but it does so by eroding meaningful human control, blurring civilian and military space, and replacing the old world of “big, slow, and certain” strategic weapons with one that is small, fast, and deeply ambiguous.

Version: 2026-03-26 22:20:16 SHA1: 1b9b65ab32

Index Terms—drone warfare, autonomous weapons, deterrence, mass disruption, infrastructure attacks, European security, algorithmic escalation, meaningful human control

References

5

I. INTRODUCTION

European security has long rested on a combination of alliance structure, conventional force posture, and the nuclear umbrella provided by the United States. If that umbrella were perceived as weakening or politically unreliable, states on NATO’s eastern and central flank would almost certainly revisit the question of sovereign deterrence [1]–[3]. In the twentieth century that question pointed toward the classical nuclear triad: land-based missiles, strategic bombers, and ballistic missile submarines. In the twenty-first century, however, the technological environment is different. Artificial intelligence, cheap sensors, distributed autonomy, and mass-produced drones invite a different line of thought [4], [5].

The most interesting and disturbing implication is that the search for deterrence may no longer focus exclusively on nuclear yield. Instead, it may focus on the ability to impose *unacceptable systemic cost*. Recent literature on strategic non-nuclear weapons suggests that emerging and disruptive technologies may create coercive effects that are strategically meaningful even if they do not replicate nuclear blast damage [6], [7]. A state does not need to vaporize a city to create a strategic disaster. If it can reliably collapse electrical distribution, telecommunications, rail transfer nodes, water pumping, and digital command systems across a country within minutes or hours, it can generate effects that are coercive on a national scale. The city remains standing, but normal life stops. Hospitals degrade, transport freezes, supply chains fail, emergency response fragments, and public authority begins to lose coherence.

This document argues that drone swarms are best understood in this context as potential instruments of *mass disruption*. They challenge strategic stability not because they perfectly reproduce the destructive logic of nuclear arms, but because they offer a cheaper, more deniable, more distributable route to nationwide paralysis. That possibility is strategically seductive and ethically corrosive at the same time. It also operates in the shadow of the nuclear taboo: states may imagine such systems as more usable precisely because they do not carry the same symbolic and normative burden as nuclear weapons, even while still threatening mass civilian suffering [7], [8].

CONTENTS

I	Introduction	1
II	From Nuclear Destruction to Distributed Disruption	2
III	Why the Logic Appeals to States	2
III-A	Affordability, Scale, and Saturation	2
III-B	Opacity and the Politics of Hidden Deterrence	3
III-C	Decision-Time Compression	3
IV	Scenario Presentations	3
IV-A	Scenario 1: Distributed Autonomous Deterrence	3
IV-B	Scenario 2: The Dark Night	3
V	Scenario: A Bad Dawn on the Eastern Flank	4
V-A	Phase 1: Detection	4
V-B	Phase 2: Automated Defense	4
V-C	Phase 3: Mutual Misreading	4
VI	Ethical and Strategic Implications	4
VI-A	Meaningful Human Control Under Pressure	4
VI-B	Civilian Space as Battlespace	4
VI-C	Accountability and Attribution	4
VII	Conclusion	5




Figure 1. Illustrative rendering of a containerized drone-swarm launch. In this paper the image functions as a scenario aid that visualizes the imagination of distributed deterrence; it is not presented as documentary evidence.

II. FROM NUCLEAR DESTRUCTION TO DISTRIBUTED DISRUPTION

The classical nuclear age was organized around a small number of highly visible, capital-intensive delivery systems. Even when terrifying, that architecture had a certain strategic legibility. Silos could be counted, submarine patrol patterns estimated, bomber bases monitored, and missile launches attributed. Stability did not arise from trust, but from mutual knowledge of where the main instruments of catastrophe were located and how quickly they could be used [1], [2].

A drone-centered deterrent points in a different direction. It is not anchored in a handful of exquisite platforms. Instead, it imagines a country-wide mesh of launch points, storage sites, autonomous vehicles, replacement stocks, and software-defined targeting logic. The promise of such a system is not elegance; it is distribution. Redundancy substitutes for survivability in the traditional sense. If enough launch assets are dispersed, concealed, or mobile, then the adversary can no longer be confident that a first strike will disarm the defender. Strategic confidence is undermined not by a single invulnerable submarine, but by the inability to map the full deterrent [5], [6].

This is where the discussion begins to move from mass destruction to mass disruption. Nuclear weapons concentrate vast destructive power in a small number of warheads. Drone swarms disperse smaller effects across many nodes, but the aggregate outcome may still be national in scope. The target set is also different. Instead of primarily threatening urban annihilation, the swarm logic threatens societal function itself:

- electrical infrastructure that supports heating, cooling, industrial production, and hospital care;
- digital communications that enable command, commerce, and civil coordination;
- transport chokepoints that move food, fuel, ammunition, and emergency services;
- symbolic and administrative nodes whose disruption signals loss of control.

The strategic message becomes clear: *we may not destroy your society in a single flash, but we can switch it off quickly enough to make organized resistance and normal governance collapse.* That is a different deterrent grammar, and possibly a more ambiguous one. The ambiguity matters because it lowers the psychological barrier to contemplating use. Political leaders who would recoil from nuclear employment may be tempted by an option that promises devastating coercion without radioactive fallout or immediate mass incineration [8], [9].

III. WHY THE LOGIC APPEALS TO STATES

A. Affordability, Scale, and Saturation

The most obvious attraction of drone-based deterrence is cost structure. Nuclear forces require advanced industrial infrastructure, fissile material, specialist delivery platforms, and the political willingness to manage a taboo-laden arsenal. Large autonomous drone inventories, by contrast, can in principle be assembled incrementally, distributed widely, and refreshed over time. They are not cheap in an absolute sense,

but they may appear far more attainable than a sovereign nuclear force [4], [7].

Scale matters because saturation changes the economics of defense. If an attacker can present a defender with many simultaneous aerial problems, the defender may be forced into a losing exchange even when its interceptors are tactically successful. The point is not that every drone survives. The point is that enough survive. At national scale, a modest leakage rate can still produce strategic-level harm [10], [11].

B. Opacity and the Politics of Hidden Deterrence

Drone deterrence also benefits from opacity. A nuclear arsenal is ultimately a declaratory and inspectable instrument: its political purpose depends on being recognized. A distributed drone arsenal can be partly visible and partly hidden. This creates a more unstable form of deterrence because uncertainty migrates from warhead counts to infrastructure ambiguity. Warehouses, transport depots, maintenance hubs, and civilian logistics spaces start to acquire strategic suspicion. The line between normal economic circulation and pre-positioned coercive capability becomes harder to draw [12], [13].

That ambiguity may appear advantageous to the defender, but it imposes a serious normative cost. Once civilian infrastructure becomes entangled with deterrence posture, the entire rear area is quietly militarized. What looks like strategic resilience from one perspective looks like the erosion of civilian immunity from another.

C. Decision-Time Compression

The deepest problem is temporal. Drone swarms do not simply threaten physical targets; they accelerate the tempo of decision. Detection, classification, routing, interception, and re-tasking increasingly become machine-speed processes. This compresses the interval in which human beings can absorb uncertainty, question sensor fusion, and resist automated escalation. A strategic order that prizes reaction time above judgment will gradually redefine human deliberation as a liability [11], [14]–[16].

This is where the thought experiment becomes most dangerous. Once both offense and defense are optimized for speed, deterrence no longer rests on slow, legible signaling. It rests on software thresholds, probabilistic classification, and brittle assumptions about what the other side's automation will interpret as hostile. The old fear of accidental launch remains, but it now takes a more distributed and more opaque form. The broader autonomous-weapons literature is helpful here: it repeatedly shows that nominal human involvement is not the same thing as meaningful human control if the time available for review has already been hollowed out by automation [17], [18].

IV. SCENARIO PRESENTATIONS

Motivation: The preceding sections describe the strategic logic of drone-based deterrence in abstract terms. The following scenarios adapt the presentation style used elsewhere in this directory in order to make the argument more concrete.

They are analytical thought experiments, not predictions and not intelligence assessments.

Scope and Restraint: The scenarios intentionally emphasize strategic effect, humanitarian consequence, and escalation logic. They do not attempt to provide operational instructions for building, deploying, or optimizing a mass drone capability.

A. Scenario 1: Distributed Autonomous Deterrence

Scenario 1: Distributed Autonomous Deterrence

1) Phase 1: The Search for a Sovereign Deterrent:

A mid-sized European state concludes that alliance guarantees can no longer be treated as automatic. It does not immediately pursue a nuclear arsenal. Instead, it constructs a distributed deterrent concept around autonomous drones, dispersed storage, and resilient replenishment.

- **Strategic Logic:** The state seeks a second-strike-like posture without relying on a small number of vulnerable launch platforms.
- **Political Message:** The deterrent claim is that no adversary can disable the system in a single blow, and no offensive plan can assume a clean disarmament.
- **Psychological Effect:** Uncertainty becomes part of deterrence. The adversary must assume that significant retaliatory capacity will survive.

2) Phase 2: Crisis Signaling Through Ambiguity:

During a regional crisis, the state does not mobilize bombers or missiles in the classical sense. Instead, it shifts logistics patterns, hardens key networks, and signals readiness through public references to national resilience and autonomous defense.

- **What Changes:** Deterrence signaling moves from visible launch platforms to diffuse indicators of survivable distributed capacity.
- **What Worsens:** Civilian infrastructure becomes entangled with strategic posture, making the rear area look increasingly like a legitimate military problem for the adversary.
- **Strategic Consequence:** Deterrence may become more survivable, but it also becomes less transparent and therefore less stable.

B. Scenario 2: The Dark Night

Scenario 2: The Dark Night

1) Phase 1: Mass Disruption Without Nuclear Use:

An aggressor launches a large autonomous swarm, not primarily against troop concentrations, but against the systems that make metropolitan life possible. The first wave targets electrical switching capacity, communications relays, rail transfer nodes, and selected fuel distribution points.

- **Immediate Effect:** The city is not destroyed phys-

ically, but it becomes progressively unworkable.

- **Operational Logic:** The goal is not spectacular blast damage; it is cascading dysfunction.
- **Public Meaning:** Citizens experience the attack as a collapse of normality rather than a single catastrophic moment.

2) *Phase 2: Coercion Through Societal Paralysis:*

Within hours, blackouts spread, digital payments fail intermittently, transport is disrupted, and emergency services struggle to coordinate. Hospitals begin triage under degraded conditions. Political leaders now face a coercive message: escalation will bring deeper paralysis.

- **Deterrent Lesson:** Drone swarms can threaten national endurance, not only battlefield positions.
- **Ethical Lesson:** An attack framed as “non-nuclear” can still produce mass suffering on a strategic scale.
- **Legal Lesson:** Once dual-use systems become the main target set, proportionality and civilian protection become extraordinarily difficult to preserve.

V. SCENARIO: A BAD DAWN ON THE EASTERN FLANK

Scenario 3: A Bad Dawn on the Eastern Flank

Content Warning: This scenario describes a fast-moving crisis in which automated systems, ambiguous attribution, and infrastructure attacks combine to produce large-scale civilian harm.

A. Phase 1: Detection

Just before dawn, border surveillance networks detect an unusual pattern of low-altitude movement and acoustic signatures. The data is incomplete. Some signals suggest a major launch event; others are noisy and contradictory. Human operators have minutes, perhaps less, to decide whether the event represents a probe, a false alarm, or the opening move of a strategic strike.

B. Phase 2: Automated Defense

Because manual review is too slow, defensive AI systems begin allocating interceptors, prioritizing critical nodes, and recommending emergency infrastructure shutdowns. These actions are technically defensive, but they are also visible. Across the border, the adversary’s systems register sudden changes in posture and begin classifying them as preparation for retaliation.

C. Phase 3: Mutual Misreading

What began as an ambiguous detection problem mutates into a reciprocal automation problem. One side sees emergency defense; the other sees imminent attack. Swarms are redirected, defenses intensify, and both governments inherit machine-accelerated facts before political leaders can establish a shared picture of what is actually happening.

The resulting damage is uneven but severe. A portion of the incoming drones are intercepted, yet enough get through to disable selected infrastructure, trigger blackouts, and provoke public panic. The attack does not produce nuclear-scale destruction, but it produces strategic consequences all the same: civilian fear, political pressure, and a strong temptation to escalate before the next wave arrives.

Ethical Analysis: This scenario illustrates how algorithmic speed can hollow out meaningful human control even when both sides believe they are acting defensively. It also shows why attribution and interpretation become central ethical variables in an age of autonomous swarms. The issue is not only whether a target classification is correct, but whether a political system still has enough time to distinguish signal from noise before violence cascades.

VI. ETHICAL AND STRATEGIC IMPLICATIONS

A. Meaningful Human Control Under Pressure

The attraction of drone deterrence is inseparable from the temptation to automate its most time-sensitive elements. Yet the more a system relies on machine-speed detection, routing, and engagement, the less plausible it becomes to claim that human beings retain meaningful control over the use of force. Human involvement can survive formally while disappearing functionally. Leaders may still set the outer parameters, but the operational decisions that matter most increasingly occur below the threshold of human comprehension and intervention [17]–[19].

B. Civilian Space as Battlespace

Drone-based mass disruption also expands the battlespace deep into civilian life. Electrical substations, telecom hubs, rail links, ports, warehouses, and data centers all become strategically salient. Some are dual-use in a narrow technical sense; many are simply indispensable to ordinary society. This means that even a doctrine aimed at coercion rather than annihilation can quickly become indiscriminate in effect. A strategy that promises bloodless systemic paralysis is often only bloodless in the imagination of planners [7], [12].

C. Accountability and Attribution

Accountability becomes harder in at least three ways. First, many consequential decisions are distributed across software layers, sensor fusion, and pre-authorized response logic. Second, swarm attacks can be structured to maximize ambiguity about origin and intent. Third, the political demand for immediate counter-action may outpace forensic certainty. Under those conditions, the temptation is to shift from confident attribution to probabilistic retaliation. That is a profoundly dangerous norm [11], [13], [19].

VII. CONCLUSION

If European states begin to think seriously about a future in which alliance guarantees cannot be taken for granted, they may search for deterrents that are cheaper, more survivable, and more scalable than a sovereign nuclear force. Drone swarms will therefore be part of the strategic imagination. Their appeal is easy to understand: they seem to offer distribution instead of vulnerability, saturation instead of singularity, and coercive leverage without nuclear use.

But that appeal rests on a hazardous bargain. Drone-based deterrence replaces the legibility of the old strategic order with opacity, dual-use entanglement, and decision-time compression. It encourages states to think in terms of nationwide systemic paralysis while normalizing the automation of crisis response. The result is not a cleaner version of deterrence. It is a more ambiguous, more civilianized, and potentially more brittle one. Recent work on AI in nuclear operations makes a similar point in a neighboring domain: the risks and benefits of AI are context-dependent, and technical sophistication does not automatically translate into strategic stability or better human control [20], [21].

The old world of strategic weapons was terrifying because it was enormous. The emerging world of autonomous swarms may be terrifying for the opposite reason: it is granular, distributed, and always on the verge of slipping beyond deliberate political control.

REFERENCES

- [1] T. C. Schelling, *Arms and Influence*. New Haven, CT: Yale University Press, 1966.
- [2] K. N. Waltz, *The Spread of Nuclear Weapons: More May Be Better*. International Institute for Strategic Studies, 1981, no. 171.
- [3] K. V. Bruusgaard, "Deterrence asymmetry and strategic stability in Europe," *Journal of Strategic Studies*, vol. 47, no. 3, pp. 334–362, 2024.
- [4] North Atlantic Treaty Organization, "Nato artificial intelligence strategy," Brussels, 2021. [Online]. Available: https://www.nato.int/cps/en/natohq/official_texts_187617.htm
- [5] J. Johnson, "Deterrence in the age of artificial intelligence & autonomy: a paradigm shift in nuclear deterrence theory and practice?" *Defense & Security Analysis*, vol. 36, no. 4, pp. 422–448, 2020.
- [6] A. Futter and B. Zala, "Strategic non-nuclear weapons and the onset of a third nuclear age," *European Journal of International Security*, vol. 6, no. 3, pp. 257–277, 2021. [Online]. Available: <https://www.cambridge.org/core/journals/european-journal-of-international-security/article/strategic-nonuclear-weapons-and-the-onset-of-a-third-nuclear-age/91EEB3B77D348252815F9F7B59DB8A32>
- [7] T. Sauer, "The potentially revolutionary impact of emerging and disruptive technologies and strategic conventional weapons on the nuclear deterrence debate," SIPRI, EU Non-Proliferation and Disarmament Paper, 2024. [Online]. Available: <https://www.sipri.org/publications/2024/eu-non-proliferation-and-disarmament-papers/potentially-revolutionary-impact-emerging-and-disruptive-technologies-and-strategic-conventional>
- [8] N. Tannenwald, "The nuclear taboo: The United States and the normative basis of nuclear non-use," *International Organization*, vol. 53, no. 3, pp. 433–468, 1999.
- [9] A. Futter, "Disruptive technologies and nuclear risks: What's new and what matters," *Survival*, vol. 64, no. 1, pp. 99–120, 2022.
- [10] M. C. Horowitz, "When speed kills: Lethal autonomous weapon systems, deterrence and stability," *Journal of Strategic Studies*, vol. 42, no. 6, pp. 764–788, 2019. [Online]. Available: <https://www.tandfonline.com/doi/full/10.1080/01402390.2019.1621174>
- [11] Y. H. Wong, J. Yurchak, R. W. Button, A. Frank, B. Laird, O. A. Osoba, R. Steeb, B. N. Harris, and S. Bae, "Deterrence in the age of thinking machines," RAND Corporation, Technical Report, 2020. [Online]. Available: https://www.rand.org/pubs/research_reports/RR2797.html
- [12] J. M. Acton, "Escalation through entanglement: How the vulnerability of command-and-control systems raises the risks of an inadvertent nuclear war," *International Security*, vol. 43, no. 1, pp. 56–99, 2018. [Online]. Available: <https://direct.mit.edu/isec/article/43/1/56/12199/Escalation-through-Entanglement-How-the>
- [13] V. Chernavskikh, "Nuclear weapons and artificial intelligence: Technological promises and practical realities," SIPRI, Background Paper, 2024. [Online]. Available: <https://www.sipri.org/publications/2024/sipri-background-papers/nuclear-weapons-and-artificial-intelligence-technological-promises-and-practical-realities>
- [14] J. R. Boyd, "The essence of winning and losing," *Air University Review*, vol. 37, no. 4, pp. 2–11, 1986. [Online]. Available: https://www.airuniversity.af.edu/Portals/10/AUPress/Books/B_0151_BOYD_DISCOURSE_WINNING_LOSING.PDF
- [15] T. Lysaug, "Hva skal vi med autonome våpen: En litteraturstudie om autonome våpensystemer," Master thesis, Forsvarets høgskole, 2021. [Online]. Available: <http://fhs.brage.unit.no/fhs-xmlui/handle/11250/2835213>
- [16] J. Johnson, "Inadvertent escalation in the age of intelligent machines: A new model for nuclear risk in the digital age," *European Journal of International Security*, vol. 7, no. 3, pp. 337–359, 2022.
- [17] R. Crotoft and D. Richemond-Barak, "The future of warfare: National positions on the governance of lethal autonomous weapons systems," 2024. [Online]. Available: <https://lieber.westpoint.edu/future-warfare-national-positions-governance-lethal-autonomous-weapons-systems/>
- [18] R. Crotoft, "What is meaningful human control, anyway? cracking the code on autonomous weapons and human judgment," 2024. [Online]. Available: <https://mwi.westpoint.edu/what-is-meaningful-human-control-anyway-cracking-the-code-on-autonomous-weapons-and-human-judgment/>
- [19] International Committee of the Red Cross, "A legal perspective: Autonomous weapon systems under international humanitarian law," 2024. [Online]. Available: <https://www.icrc.org/en/document/autonomous-weapon-systems-under-international-humanitarian-law>
- [20] H. S. Lin, "Artificial intelligence and nuclear weapons: A commonsense approach to understanding costs and benefits," *Texas National Security Review*, vol. 8, no. 3, pp. 98–109, 2025. [Online]. Available: <https://tnsr.org/2025/06/artificial-intelligence-and-nuclear-weapons-a-commonsense-approach-to-understanding-costs-and-benefits/>
- [21] T. McDonnell, M. Chesnut, T. Ditter, A. Fink, and L. Lewis, "Artificial intelligence in nuclear operations: Challenges, opportunities, and impacts," CNA, Research Memorandum, 2023. [Online]. Available: <https://www.cna.org/reports/2023/04/Artificial-Intelligence-in-Nuclear-Operations.pdf>